AI2100/AI5100 Deep Learning, Spring 2022

Indian Institute of Technology Hyderabad

Homework 3, Convolutional Neural Networks (CNNs), Assigned 06.03.2022, Due **11:59 pm on 16.03.2022**

Ask the right questions, and nature will open the doors to her secrets. – Sir C. V. Raman

Instructions:

- It is **strongly recommended** that you work on your homework on an *individual* basis. If you have any questions or concerns, feel free to talk to the instructor or the TAs.

- **The functions implemented for HW2 can be reused here.**

- Use `matplotlib` where specified - `https://matplotlib.org/tutorials/introductory/images.html`.

- Use color images from the CIFAR-10 dataset `https://www.cs.toronto.edu/~kriz/cifar.html`.

- Please turn in Python Notebooks with the following notation for the file name: `your-roll-number-hw3.ipynb`.

- Do not turn in images. Please use the same names for images in your code as in the database. The TAs will use these images to test your code.

In this assignment you will implement each of the components of a Convolutional Neural Network (CNN) from scratch (i.e., without using built-in functions for convolution, pooling, non-linearity, padding, striding). Your implementation must accept an image input and generate an output vector. Use random weights for filter kernels and fully connected layers. Specifically, implement the following:

1. **Convolution function:** It accepts as input an image, a filter kernel, stride, padding and the non-linear function. The function must convolve the input image (after padding if specified) with the kernel (at the specified stride size) and generate an output activation after applying the specified non-linearity. Verify with the standard options for the non-linear activation functions - `sigmoid`, `tanh`, `ReLU`, `Parametric ReLU (PReLU)`. Display the input image, the filter kernel and the output activation map. Ensure that your function can accept multi-channel input and a corresponding kernel volume. (3)

2. **Pooling function:** It accepts as input the activation map output from the convolution function, a pooling function, and stride. The function must output the appropriately pooled activation map. Display the input activation map and the pooled output. (2)

3. **Convolution layer function:** It accepts as input a volume (image or activation maps), filter kernels, stride, padding and the non-linear function. The function must convolve the input volume (after padding if specified) with each of the kernels (at the specified stride size) and generates an output activation volume after applying the specified non-linearity. Display the input image or activation maps, the filter kernels and the output activation maps. Verify that the output of this function does indeed have the expected size ($W \times H \times C$) as discussed in class. (1)

4. **Pooling layer function:** It accepts as input the activation map volume, the pooling function, stride, and generates a pooled output volume. Display the input and output volumes. (1)

5. **Flattening (unraveling) function:** It accepts as input the activation map volume output by the pooling layer and generates a vector of a specified size. It is important to note that this function has a weight matrix associated with it whose size is chosen such that the input and desired output sizes are matched. (1)

6. **Multilayer Perceptron (MLP) function (Fully Connected):** It accepts as input a vector, the number of hidden layers, the size of each hidden layer, the non-linear function, and the size of the output layer. This function should generate an output vector of the specified size. Generate the output with and without the *softmax* function applied to the output layer. (2)

7. **Feed-forward path:** Finally, use the functions you have written to implement a CNN with the following architecture. The CNN must accept an image input and output a vector of appropriate dimension. In other words, the function must effectively implement the feed-forward path in a CNN. (3)

- Input image of size $32 \times 32 \times 3$. Use images from the CIFAR-10 dataset.
- Convolution layer with 16 kernels of size $3 \times 3 \times 3$ and sigmoid activation.
- Max pooling layer of size 2×2 with a stride of 2 along each dimension.
- Convolution layer with 8 kernels of size $3 \times 3 \times C_2$ and sigmoid activation. What is the value of C_2?.
- Max pooling layer of size 2×2 with a stride of 2 along each dimension.
- A flattening layer. Note that the input size is fixed by the previous layer's output. Your flattening matrix can be square, tall, or fat. Any choice is fine for this problem. Pick any one.
- An MLP with one hidden layer that accepts as input the flattening layer output and maps it to 10 output nodes. Use sigmoid activation for the MLP.

Verify that your composition of function accepts and image input and outputs a vector.

8. (a) Choose an image from each of the 10 classes and display the output vector for each case. Do you see any trend in the output vectors? (1)

   (b) Does a randomly initialized network show any discriminability? Visualize the bottleneck layer (output of flattening layer) using *builtin* tSNE plots. Choose three images per class. If you are wondering about this question, check out the Deep Image Prior paper. (1)